

drafter source	drafter training							target conditioning				draft shape								verification					memory				execution					
	training-free	trained on fixed data	distilled from target	trained on own drafts	built into target training	acceptance-aligned objective	online / test-time adaptation	multi-layer target features	reuses rejected KV	re-anchors at rejections	single chain	static tree	dynamic / optimised tree	path-conditioned tree	lattice / path search	multiple drafts	adaptive length	long block (>= 24)	greedy match only	token-wise rejection	multi-draft coupling	block / sequence-level	cross-vocabulary	sparse / streaming KV	quantised weights / KV	constant-memory drafter	offloaded target	hierarchy of drafters	overlap draft & verify	drafter on other devices	load-aware speculation	latency-SLO-aware	RL-rollout aware	
separate small LM	3	1	3			1	1					10	5	2			3	3			15	4	2	1	2			2	3	4	3	1		
target's own layers / weights	5	1			1							5					1			9				3	2			1			1			
parallel heads on target		2			1							2	1						1	2														
AR head on target features	1	4	3	6		2	2	1	1			2	3	11				1	1	15		1				1				1		2		
lookup: prompt / references	5											3	1	1			2		1	5										1				
lookup: datastore / history	6											1	2	1	1		2			6												1		
lookup: sibling samples	5													2						5												3		
Jacobi n-gram pool	1															1				1														
parallel / diffusion drafter	8	2	9	3		4	2	5			1	11		6	3	1	1	8	1	10	15	1	1							5				
conditioned parallel drafter	4	1	9	1		3		9		1		8			7	2		2		4	12	1	1							3				

Number of papers whose drafter is of this kind and that also have the column's trait (● or ○ in the matrix). Dashed empty cell: no paper tried it yet (not every one makes sense).