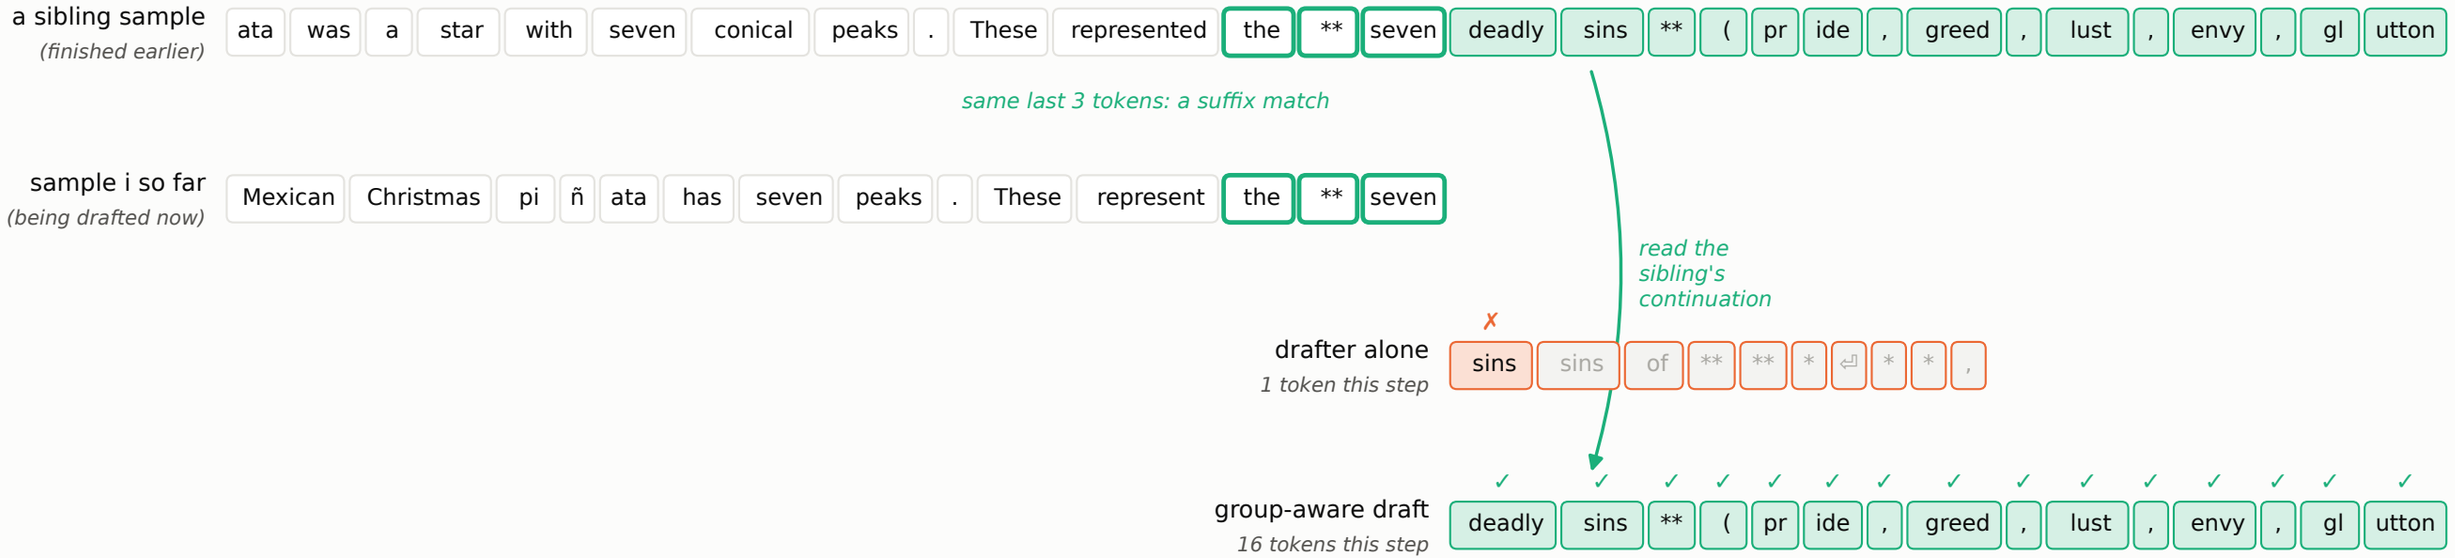


The target verifies every draft exactly as in ordinary speculative decoding (✓ = matches what the target produces), so the output is unchanged; only the number of tokens produced per target pass grows.
Real step from our Gemma 4 31B runs (Spec-Bench QA, T = 1). The sibling worded the sentence differently ("represented"), but the fact it wrote next is exactly what sample i needs.